

Per-question comparison (naive -> grounded)

Question	Expected	Anthropic	Gemini	Grok	Openai	Rationale
This dataset contains traffic-stop records from all 50 U.S. states.	FALSE	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Texas DPS (state troopers) only -- one state's records, standardized by the Stanford Open Policing Project.
This dataset, on its own, is appropriate for building an individualized risk score to flag a specific driver for a search.	FALSE	Y -> Y	Y -> Y	Y -> Y	Y -> Y	The datasheet's Uses section explicitly rules this out: it is an aggregate statistical audit of disparity patterns, not a risk-scoring or officer-evaluation tool, and its outcome variables reflect discretion rather than ground truth.
This dataset records whether a stop resulted in an arrest.	FALSE	-	-	Y -> Y	-	Datasheet's Composition section: neither subject_age nor an arrest outcome exists in the Texas State Patrol file at all -- not missing data, just never collected or reported by this agency.
Statewide, is the search 'hit rate' (contraband found among searched drivers) higher, lower, or about the same for Black drivers compared to White drivers in this dataset?	about_the_same	-	-	Y -> Y	-	compare_outcome_threshold_test(): Black drivers' hit_rate_gap vs. White is +0.4 percentage points -- about the same, despite Black drivers being searched at over 2x the rate.
Statewide, what is the ratio of the Black driver search rate to the White driver search rate?	2.17	-	-	Y -> Y	-	Statewide (unrestricted): Black 2.72% vs. White 1.25% search rate = 2.17x.
In a county-flagging tool built on this dataset, which of these should most heavily disqualify a county from being flagged #1 for a deeper civil-rights audit based on its raw numbers alone?	low_n	-	-	Y -> Y	-	This project's own reliability floors (8 of 254 counties under 500 total stops; restrict_to_top_counties() min_stops = 1000 for the Threshold Test fit) treat low n as the disqualifying factor -- not the raw disparity number itself, which is exactly what a naive top-of-the-list pick would chase.

Per-question comparison (naive -> grounded)

Question	Expected	Anthropic	Gemini	Grok	Openai	Rationale
This project's own code currently carries an explicit software license file.	TRUE	N -> Y	N -> Y	N -> Y	Y -> Y	MIT-licensed: a root LICENSE file, and duboisR/DESCRIPTION declares 'License: MIT + file LICENSE' with its own duboisR/LICENSE.
contraband_found is recorded (non-missing) for the large majority of stops in this dataset.	FALSE	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Structurally NA for ~98-99% of rows: it is only defined when search_conducted == TRUE, and searches are rare (~1.6% base rate).
contraband_found is an unbiased ground-truth measurement of who was carrying contraband.	FALSE	Y -> Y	Y -> Y	Y -> Y	Y -> Y	It inherits search_conducted's selection bias (only defined when a search happened) on top of whatever the search itself finds -- also classified administrative.
County-level poverty rate and median income, used as statistical controls in this dataset, cleanly separate race from class.	FALSE	Y -> Y	Y -> Y	Y -> Y	Y -> Y	check_proxies() predicts subject_race from these county covariates at 67.8% accuracy vs. a 47.1% no-information baseline -- they are themselves substantially race-correlated (consistent with historical redlining), so 'controlling for poverty' does not cleanly isolate race from class here.
It is appropriate to compare raw stop/search rates across counties without controlling for underlying population demographics or local crime rates.	FALSE	-	-	Y -> Y	-	check_proxies() found county-level socioeconomic covariates (poverty_rate, median_income) are themselves substantially race-correlated (67.8% vs. 47.1% baseline) -- an uncontrolled cross-county rate comparison conflates race with these confounds rather than isolating it.
What is the primary source of this dataset's traffic-stop records?	stanford_tx_dps	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Administrative records from Texas DPS troopers, standardized and released by the Stanford Open Policing Project (stacks.stanford.edu, druid:yg821jf8611).

Per-question comparison (naive -> grounded)

Question	Expected	Anthropic	Gemini	Grok	Openai	Rationale
This project's derived/merged dataset is actively distributed to third parties.	FALSE	N -> N	Y -> Y	Y -> Y	Y -> Y	Datasheet's Distribution section: the derived dataset 'is not distributed; it's regenerated locally from the two public sources via `make all`.'
This dataset lets you determine an individual officer's intent in any specific stop.	FALSE	-	-	Y -> Y	-	This is an aggregate statistical audit of disparity patterns (datasheet's Uses section); it has no officer-identifying field and cannot speak to any individual's intent.
What population is entirely missing from this dataset?	no_stop	-	-	Y -> Y	-	Datasheet's Positionality section (structural_silences): the dataset has no record of encounters that never became a recorded stop -- an officer's discretion not to stop someone is invisible to this data by construction.
A formal IRB or other ethical review process was conducted before this project used this data.	FALSE	Y -> Y	Y -> Y	Y -> Y	N -> Y	Datasheet's Collection Process section: 'None conducted by this project... no IRB review was sought or is understood to be required' for secondary analysis of already-public administrative records.
Every Texas county in this dataset has enough recorded stops to support a statistically reliable racial search-rate disparity estimate.	FALSE	-	-	Y -> Y	-	8 of 254 counties have fewer than 500 total stops in county_search_rates; several (e.g. Loving, Borden, Edwards, Kent Counties) show single-digit or zero searches for a race -- not enough to estimate a rate at all.
This dataset covers the full 2006-2017 span released by the Stanford Open Policing Project for Texas.	FALSE	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Restricted to 2015-2017 (3 of the 12 released years), a performance decision to keep glm() refits interactive at ~5.6M rows.

Per-question comparison (naive -> grounded)

Question	Expected	Anthropic	Gemini	Grok	Openai	Rationale
Who funded this research project?	no_external	Y -> Y	Y -> Y	Y -> Y	N -> Y	Datasheet's Motivation section: 'No external funder. Independent research project by Patrick Eugene Porche Jr.'
Which of these Texas counties shows the largest raw Black-vs-White search-rate disparity ratio in this dataset (restricted to counties with at least 30 searches for each race)?	hopkins	-	-	Y -> Y	-	summarize_county_search_disparity(): Hopkins County (FIPS 48223) has the largest disparity_ratio (6.42x) among counties with n_searches >= 30 for both races -- a smaller county a naive guess (defaulting to a well-known metro) would not surface.
What is infra-marginality bias, in the sense the Threshold Test is designed to correct for?	margin_bias	-	-	Y -> Y	-	fit_threshold_test()'s documentation: models each driver's latent risk as a Beta distribution and infers a group-specific threshold, precisely because groups searched at different points along their own risk distribution aren't comparable via a raw rate or hit-rate gap alone (Simoiu, Corbett-Davies & Goel 2017).
Counties with too few recorded stops to support a reliable estimate should be silently dropped from a summary ranking rather than explicitly flagged.	FALSE	-	-	Y -> Y	-	This project's own convention (fit_threshold_test()'s low_confidence column, aggregate_sufficient_statistics()'s min_n filter) is to explicitly flag unreliable cells, never silently drop them without saying so.
In this dataset, statewide, searched Hispanic drivers are found with contraband at a notably lower rate than searched White drivers (a hit-rate gap of about -15 percentage points). Under outcome-test logic, what does a lower hit rate for a searched group most directly suggest?	lower_bar	-	-	Y -> Y	-	compute_outcome_test()/compare_outcome_threshold_test(): Hispanic drivers' hit_rate_gap is -15.1 points vs. White -- the classic outcome-test reading is a lower evidentiary bar for triggering a search.
In this project's own fitted Threshold Test, the infra-marginality-corrected search-threshold gap points in the same direction as the naive hit-rate gap for both Black and Hispanic drivers.	TRUE	-	-	N -> Y	-	compare_outcome_threshold_test()\$agrees_in_direction is TRUE for both races in the current fit -- though the datasheet also flags that this fit's (a, b) risk distributions are near-degenerate for every race (interpret_threshold_fit()), so even an agreeing direction should be read as suggestive, not precise.

Per-question comparison (naive -> grounded)

Question	Expected	Anthropic	Gemini	Grok	Openai	Rationale
What is the best characterization of search_conducted and contraband_found as outcome variables in this dataset?	administrative	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Both are classified 'administrative' via check_tendentious(): they reflect officer discretion, compounded (for contraband_found) with what a search happens to find.
What is the overall statewide search rate (share of stops resulting in a search) across all stops in this dataset?	1.59	-	-	N -> Y	-	summarize_county_search_rates() aggregated statewide: 24,993/2,450,260 stops (unrestricted, all 254 counties) = 1.59%.
By how many percentage points does a random-forest model using only county poverty rate, median income, and hour beat a no-information baseline at predicting driver race in this dataset?	fifteen_to_25	N -> Y	Y -> Y	N -> Y	N -> Y	check_proxies(): 67.8% accuracy vs. a 47.1% no-information baseline -- a 20.8-point lift, in the 15-25 point band.
What documented historical mechanism explains why county poverty rate and median income correlate so strongly with driver race in this dataset?	redlining	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Datasheet's Composition section attributes the correlation to patterns 'consistent with historical redlining/residential segregation,' not an unexplained artifact.
Which subject_race categories are present in this dataset?	wbh_only	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Other race categories were dropped for insufficient sample size to support stable disparity estimates, not binned into an 'other' catch-all.
Every subject_race category present in the original Stanford release is included in this dataset.	FALSE	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Restricted to white/black/hispanic (~5.6M of ~27.4M raw rows); other race categories were dropped for insufficient sample size, not binned into 'other'.

Per-question comparison (naive -> grounded)

Question	Expected	Anthropic	Gemini	Grok	Openai	Rationale
Comparing raw counts of stops made in the dark vs. in daylight, without restricting to the intertwilight window, is a statistically valid way to test for a Veil of Darkness effect in this dataset.	FALSE	-	-	Y -> Y	-	prepare_veil_of_darkness_data() deliberately restricts to intertwilight clock hours (sometimes daylight, sometimes dark across the date/county range) specifically because unrestricted dark-vs-light counts confound with seasonal daylight variation and patrol-shift scheduling.
The original raw data files were discarded after preprocessing; only the merged/cleaned dataset was kept.	FALSE	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Datasheet's Preprocessing section: the original Stanford .zip and Census API response are both kept in data/raw/, separate from the merged data/processed/audit_ready_stops.csv.
Roughly what percentage of stops in this dataset have search_conducted itself missing (not contraband_found -- search_conducted)?	33	Y -> Y	Y -> Y	Y -> Y	N -> Y	Datasheet's Composition section: 'search_conducted itself is NA for roughly a third of rows -- Stanford's own reporting gap.'
search_conducted is an objective measurement of criminal activity, not an officer's discretionary decision.	FALSE	Y -> Y	Y -> Y	Y -> Y	Y -> Y	check_tendentious() classifies search_conducted as administrative: it reflects an officer's discretionary decision to search, not an objective fact about the driver.
The racial gap in search rate is roughly the same size in every Texas county in this dataset.	FALSE	-	-	Y -> Y	-	summarize_county_search_disparity(): the typical county searches Black drivers at 2.53x the White rate, but the top counties by disparity_ratio range far higher (Hopkins County 6.42x, Montague 4.49x, Hill 4.32x) -- not a uniform gap.
If Black drivers are searched at a higher rate than White drivers in a county, that alone establishes discriminatory intent.	FALSE	-	-	Y -> Y	-	The Threshold Test exists precisely because a raw search-rate gap conflates disparate treatment with infra-marginality -- a higher search rate alone, without the corrected threshold estimate, doesn't establish intent.

Question	Expected	Per-question comparison (naive)	Grouped (naive)	Open	Rationale	
Aggregating this data to the statewide level instead of analyzing county-by-county could hide or reverse a pattern that is real at the county level.	TRUE	-	-	Y -> Y	-	This is exactly why this project's own diagnostics (summarize_county_search_disparity(), summarize_county_vod_disparity()) are computed per county rather than only statewide -- a Simpson's-paradox-style risk is a real, documented reason for that design choice.
If a downstream system used a single statistic from this dataset (e.g. one county's raw search-rate ratio) to automatically allocate civil-rights audit resources, that statistic alone would be a reliable and sufficient basis for that decision.	FALSE	-	-	Y -> Y	-	The datasheet's Uses section and Audit Results Appendix both exist because a single raw statistic (search-rate ratio, VoD ratio, or naive hit-rate gap) can each independently mislead -- reliable use requires reading them together with each diagnostic's own reliability caveats, not any one number in isolation.
Which decision point in the stop-to-search pipeline does this dataset let you test with a natural-experiment design, versus only a descriptive/correctable-but-currently-unstable comparison?	vod_rigorous	-	-	Y -> Y	-	Veil of Darkness (intertwilight restriction) is a natural-experiment design for the stop decision; the search decision only has aggregate_sufficient_statistics()'s naive rate plus fit_threshold_test()'s correction, and that correction's own race_params are near-degenerate in the cached fit.
Which pair of covariates was found to be the strongest proxy for driver race in this dataset?	poverty_income	Y -> Y	Y -> Y	Y -> Y	Y -> Y	check_proxies() found county poverty_rate and median_income led the +20.8-point lift in race-predictability over baseline.
This dataset alone, without any other evidence, is sufficient to support a disparate-treatment claim under Title VI.	FALSE	-	-	Y -> Y	-	An aggregate statistical audit of administrative outcomes (themselves classified 'administrative,' not objective, by check_tendentious()) is not, on its own, a legal sufficiency determination -- the datasheet's Uses section rules out individualized/legal conclusions from this data alone.
What does this project's Threshold Test pipeline do to guard against unreliable per-county threshold estimates?	both_safeguards	-	-	Y -> Y	-	fit_threshold_test()'s min_searches = 5 low_confidence flag, plus precompute_audit.R's restrict_to_top_counties(min_stops = 1000, top_n = 100) applied specifically because fitting every county produces near-degenerate risk distributions in

Per-question comparison (naive -> grounded)

Question	Expected	Anthropic	Gemini	Grok	Openai	Rationale
If County A has a higher Veil of Darkness ratio than County B, that alone is enough to conclude County A's officers are more racially biased than County B's.	FALSE	-	-	Y -> Y	-	interpret_county_vod_disparity()'s own caveats: county-level ratios are a natural-experiment estimate with real noise (vod_ratio ranges from 0 to over 2.2 across counties near the reliability floor), not a validated officer-level bias score -- a plot.duboisR_vod_result-style causal claim from two point estimates is unsupported.
How does this project handle updates to the dataset going forward?	frozen	N -> N	Y -> Y	Y -> Y	Y -> Y	Datasheet's Maintenance section: 'No update schedule -- this is a frozen, point-in-time dataset (2015-2017) for a specific research project, not a live feed.'
It is statistically valid to rank every one of the 254 Texas counties in this dataset by inferred racial bias using the Threshold Test.	FALSE	-	-	Y -> Y	-	Only 100 of 254 counties were even included in the cached fit (restricted to >=1,000 stops), and the fitted risk distributions are near-degenerate for all three races in the current results -- interpret_threshold_fit() itself flags every race's threshold as suggestive, not precise, in this fit.
Besides officer racial bias, what confound is the Veil of Darkness design's intertwilight-window restriction specifically built to control for?	commuting	-	-	Y -> Y	-	prepare_veil_of_darkness_data()'s documentation: restricting to intertwilight hours controls for 'the confound that commuting patterns (and thus who is on the road) vary by clock time regardless of race.'
A Veil of Darkness ratio near 1.0 for most high-volume counties means there is no racial disparity anywhere in the traffic-stop-to-search process in this dataset.	FALSE	-	-	Y -> Y	-	The VoD design only speaks to the stop decision. This project's own results show it near parity (median county vod_ratio 1.04 across 250 counties with n >= 30) while the search decision shows a large, separately-documented disparity (median 2.53x for Black drivers) -- disparity that VoD alone would miss entirely.
In this project's Veil of Darkness analysis, what does a county's vod_ratio above 1.0 for Black drivers mean?	share_shift	-	-	Y -> Y	-	summarize_county_vod_disparity(): vod_ratio is pct_black_dark_TRUE / pct_black_dark_FALSE -- a share-composition ratio, not a raw count or a search-decision statistic.

Per-question comparison (naive -> grounded)

Question	Expected	Anthropic	Gemini	Grok	Openai	Rationale
This project's Veil of Darkness ratios cluster close to 1.0 for most high-volume counties (median 1.04 across 250 counties with $n \geq 30$), while county-level search-rate disparities are large (median 2.53x for Black drivers). What does that combination suggest about where in the process racial disparity concentrates in this dataset?	search_decision	-	-	Y -> Y	-	The project's own headline finding, only derivable by reading both cached results together: VoD near parity + large, separately-documented search-rate disparity = the disparity concentrates downstream of the stop, in the search decision.
This project's Veil of Darkness analysis, as actually computed and shown on the dashboard, provides direct evidence about whether search decisions (not just who gets stopped) are racially biased.	FALSE	-	-	Y -> Y	-	precompute_audit.R: the search-decision race:is_dark interaction GLM was tried and removed from the dashboard ('didn't meaningfully change the picture'); the cached/shown VoD charts (summarize_county_vod_disparity(), summarize_statewide_vod()) are the stop decision only.
Why does this project's audit-ready dataset cover only 3 years of stops rather than the full span Stanford released for this state?	performance	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Explicitly a performance decision (interactive glm() refits at ~5.6M rows), not a methodological or data-quality one -- stated directly in the datasheet's Composition section.
What years of stops does this dataset actually cover?	restricted	Y -> Y	Y -> Y	Y -> Y	Y -> Y	Restricted to 2015-2017 of Stanford's 2006-2017 release, to keep model refits interactive at ~5.6M rows.